

Resolving Network Errors in a Clustered Environment

Summary

Network errors in a clustered environment can be caused by misconfigurations, firewall issues, DNS failures, load balancer errors, or service-to-service communication problems. This document provides **step-by-step troubleshooting** guidelines to identify, diagnose, and resolve such network issues efficiently.

Steps to Resolve Network Errors

1 Check Network Connectivity & Node Availability

Command to Run:

```
kubectl get nodes -o wide  
ping <node-ip>
```

Expected Output:

- All nodes should be in a `Ready` state.
- Ping should return a response from the destination node.

Actions If Issue Exists:

- If a node is `NotReady`, check `kubectl describe node <node-name>`.
 - If ping fails, verify network routes and firewall settings.
-

2 Verify Pod & Service Communication

◆ Command to Check Pods & Services:

```
kubectl get pods -o wide  
kubectl get svc -o wide
```

◆ Expected Output:

- Pods should be running and in `Running` state.
- Services should be exposing the correct ports.

◆ Actions If Issue Exists:

- Use `kubectl logs <pod-name>` to check for errors.
- Restart affected pods: `kubectl delete pod <pod-name>` .

3 Test Internal Cluster DNS Resolution

◆ Command to Run DNS Check:

```
kubectl exec -it <pod-name> -- nslookup <service-name>
```

◆ Expected Output:

- DNS should resolve to the correct service IP.

◆ Actions If Issue Exists:

- Restart CoreDNS: `kubectl rollout restart deployment coredns -n kube-system` .
- Check logs for DNS failures: `kubectl logs -n kube-system -l k8s-app=kube-dns` .

4 Check Network Policies & Firewalls

◆ Command to List Network Policies:

```
kubectl get networkpolicy -A
```

◆ Expected Output:

- Policies should allow traffic between necessary services.

◆ Actions If Issue Exists:

- Modify network policies to allow required communication.
 - Check firewall rules on cloud provider (AWS Security Groups, Azure NSG, GCP Firewall).
-

5 Validate Load Balancer & Ingress Configuration

◆ Command to Check Load Balancer & Ingress:

```
kubectl get ingress -o wide  
kubectl get services -o wide | grep LoadBalancer
```

◆ Expected Output:

- LoadBalancer should have an external IP assigned.
- Ingress should route traffic correctly.

◆ Actions If Issue Exists:

- Verify ingress controller logs: `kubectl logs -n ingress-nginx -l app.kubernetes.io/name=ingress-nginx`.
 - Check cloud provider's load balancer status.
-

6 Debug with Network Tracing & Logs

◆ Commands to Trace Packets:

```
kubectl exec -it <pod-name> -- tcpdump -i eth0 port <service-port>  
kubectl logs <pod-name> --previous
```

◆ Expected Output:

- TCP traffic should flow as expected.
- Logs should not contain `connection refused` or `timeout` errors.

◆ Actions If Issue Exists:

- Use `kubectl port-forward` to test service access.
- Increase logging verbosity for better insights.

✓ Conclusion

By following these steps, you can systematically **identify and resolve network issues in a Kubernetes or clustered environment**. Proper monitoring, logging, and security configurations help prevent recurring issues.

For persistent problems, consider **network observability tools like Istio, Cilium, or Calico** to gain deeper insights into network traffic.